

Import the required libraries

In [ ]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Loading of dataset

In [ ]:

```
df = pd.read_csv("/content/abalone.csv")
df
```

Out[ ]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows × 9 columns

In [ ]:

*#modifying the dataset*

df['Age'] = df.Rings + 1.5

```
df=df.rename(columns = {'Whole weight':'Whole_weight','Shucked weight': 'Shucked_weight',
                        'Viscera weight': 'Viscera_weight',
                        'Shell weight': 'Shell_weight'})
```

df

Out[ ]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	
...	...	...	...	...	...	...	...	...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	

4177 rows × 10 columns



In [ ]:

df.shape

Out[ ]:

(4177, 10)

In [ ]:

df.mean

Out[ ]:

```

<bound method NDFrame._add_numeric_operations.<locals>.mean of
ength Diameter Height Whole_weight Shucked_weight \      Sex  L
0      M    0.455    0.365    0.095    0.5140    0.2245
1      M    0.350    0.265    0.090    0.2255    0.0995
2      F    0.530    0.420    0.135    0.6770    0.2565
3      M    0.440    0.365    0.125    0.5160    0.2155
4      I    0.330    0.255    0.080    0.2050    0.0895
... ..
4172   F    0.565    0.450    0.165    0.8870    0.3700
4173   M    0.590    0.440    0.135    0.9660    0.4390
4174   M    0.600    0.475    0.205    1.1760    0.5255
4175   F    0.625    0.485    0.150    1.0945    0.5310
4176   M    0.710    0.555    0.195    1.9485    0.9455

Viscera_weight Shell_weight Rings  Age
0      0.1010    0.1500    15  16.5
1      0.0485    0.0700    7   8.5
2      0.1415    0.2100    9  10.5
3      0.1140    0.1550   10  11.5
4      0.0395    0.0550    7   8.5
... ..
4172    0.2390    0.2490   11  12.5
4173    0.2145    0.2605   10  11.5
4174    0.2875    0.3080    9  10.5
4175    0.2610    0.2960   10  11.5
4176    0.3765    0.4950   12  13.5

```

[4177 rows x 10 columns]&gt;


## Visualization

In [ ]:

```
#Univariate analysis  
plt.hist(df['Age'])
```

Out[ ]:

```
(array([ 17., 431., 1648., 1388., 329., 228., 100., 29., 4.,  
        3.]),  
 array([ 2.5, 5.3, 8.1, 10.9, 13.7, 16.5, 19.3, 22.1, 24.9, 27.7, 30.  
        5]),  
 <a list of 10 Patch objects>)
```



In [ ]:


```
sns.countplot(df['Age'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[ ]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa5afa97750>




In [ ]:

```
plt.pie(df['Sex'].value_counts(),labels=['Male','Female','Infant'])
```

Out[ ]:

```
([<matplotlib.patches.Wedge at 0x7fa5af4df050>,  
 <matplotlib.patches.Wedge at 0x7fa5af560310>,  
 <matplotlib.patches.Wedge at 0x7fa5af4df910>],  
 [Text(0.45010440780275796, 1.0036961801643607, 'Male'),  
  Text(-1.0848393519507589, -0.18199884741134378, 'Female'),  
  Text(0.6099659291018239, -0.9153914820091724, 'Infant')])
```




In [ ]:

```
#Bi-varient analysis:  
sns.barplot(x = df.Sex,y = df.Whole_weight)
```

Out[ ]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa5af4f8890>

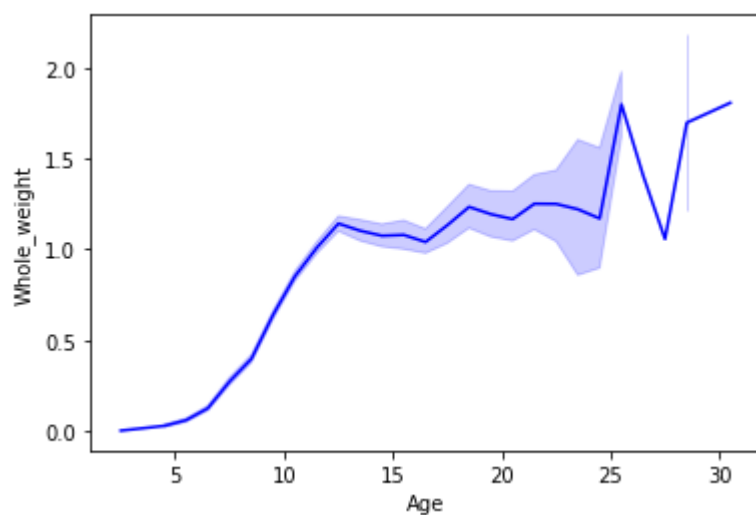


In [ ]:

```
sns.lineplot(x = df.Age, y = df.Whole_weight, color = 'blue')
```


Out[ ]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa5af458590>



In [ ]:

```
plt.scatter(df.Shucked_weight,df.Shell_weight)
plt.xlabel("Shucked_weight")
plt.ylabel("Shell_weight")
sns.set_style('ticks')
```



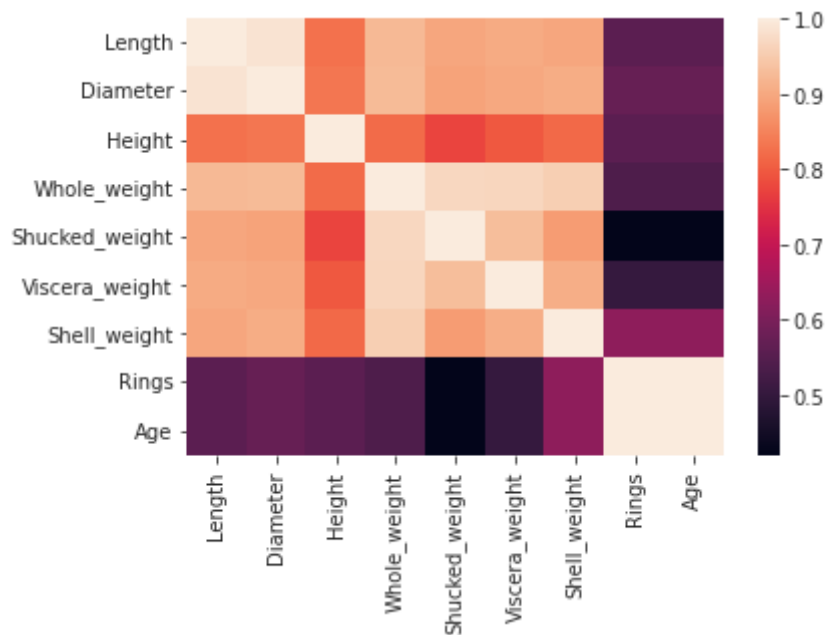
In [ ]:

```
#multi variant analysis

corr = df.corr()
sns.heatmap(corr,xticklabels=corr.columns,yticklabels=corr.columns)
```

Out[ ]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa5af36fe50>





```
sns.pairplot(df)
plt.show()
```



In [ ]:

```
df.describe()
```

Out[ ]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weigh
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.18059
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.10961
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.00050
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.09350
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.17100
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.25300
<b>max</b>	0.815000	0.650000	1.130000	2.825500	1.488000	0.76000

## Handling missing values

In [ ]:

```
df.isnull().sum()
```

Out[ ]:

```
Sex          0
Length       0
Diameter     0
Height       0
Whole_weight 0
Shucked_weight 0
Viscera_weight 0
Shell_weight 0
Rings        0
Age          0
dtype: int64
```

## Outliers handling

In [ ]:

```
outliers=df.quantile(q=(0.25,0.75))
outliers
```

Out[ ]:

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
<b>0.25</b>	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130
<b>0.75</b>	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329

In [ ]:

```
q1 = df.Age.quantile(0.25)
q2 = df.Age.quantile(0.75)
q3 = q2 - q1
lower_limit = q1 - 1.5 * q3
df.median(numeric_only=True)
```

Out[ ]:


```
Length      0.5450
Diameter    0.4250
Height      0.1400
Whole_weight 0.7995
Shucked_weight 0.3360
Viscera_weight 0.1710
Shell_weight 0.2340
Rings       9.0000
Age        10.5000
dtype: float64
```

In [ ]:

```
df['Age'] = np.where(df['Age'] < lower_limit, 7, df['Age'])
sns.boxplot(x=df.Age, showfliers = False)
```

Out[ ]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa5ab05f090>



## Encoding

In [ ]:

```
from sklearn.preprocessing import LabelEncoder
encode = LabelEncoder()
df.Sex = encode.fit_transform(df.Sex)
df
```

Out[ ]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	
...	...	...	...	...	...	...	...	...
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390	
4173	2	0.590	0.440	0.135	0.9660	0.4390	0.2145	
4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875	
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765	

4177 rows × 10 columns



Independent and Dependent Variables

In [ ]:

```
#independent variable
```

```
x = df.iloc[:,8]  
x
```

Out[ ]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	
...	...	...	...	...	...	...	...	...
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390	
4173	2	0.590	0.440	0.135	0.9660	0.4390	0.2145	
4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875	
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765	

4177 rows × 8 columns



In [ ]:

*#dependent variable*

```
y = df.iloc[:,9:]
y
```

Out[ ]:

	Age
0	16.5
1	8.5
2	10.5
3	11.5
4	8.5
...	...
4172	12.5
4173	11.5
4174	10.5
4175	11.5
4176	13.5

4177 rows × 1 columns

## Scaling

In [ ]:

```
from sklearn import preprocessing
standardisation = preprocessing.StandardScaler()
new_x = standardisation.fit_transform(x)
print(new_x)
```

```
[[ 1.15198011 -0.57455813 -0.43214879 ... -0.60768536 -0.72621157
  -0.63821689]
 [ 1.15198011 -1.44898585 -1.439929    ... -1.17090984 -1.20522124
  -1.21298732]
 [-1.28068972  0.05003309  0.12213032 ... -0.4634999  -0.35668983
  -0.20713907]
 ...
 [ 1.15198011  0.6329849   0.67640943 ...  0.74855917  0.97541324
  0.49695471]
 [-1.28068972  0.84118198  0.77718745 ...  0.77334105  0.73362741
  0.41073914]
 [ 1.15198011  1.54905203  1.48263359 ...  2.64099341  1.78744868
  1.84048058]]
```

In [ ]:

```
min_max_scaler = preprocessing.MinMaxScaler()
new_x = min_max_scaler.fit_transform(x)
print(new_x)
```

```
[[1.          0.51351351 0.5210084  ... 0.15030262 0.1323239  0.14798206]
 [1.          0.37162162 0.35294118  ... 0.06624075 0.06319947 0.06826109]
 [0.          0.61486486 0.61344538  ... 0.17182246 0.18564845 0.2077728 ]
 ...
 [1.          0.70945946 0.70588235  ... 0.3527236  0.37788018 0.30543099]
 [0.          0.74324324 0.72268908  ... 0.35642233 0.34298881 0.29347285]
 [1.          0.85810811 0.84033613  ... 0.63517149 0.49506254 0.49177877]]
```

## Split the data into training and testing

In [ ]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y)
```

## Model building

In [ ]:

```
from sklearn.linear_model import LinearRegression
mlr=LinearRegression()
mlr.fit(x_train,y_train)
```

Out[ ]:

LinearRegression()

## Training and testing

In [ ]:

```
x_test
```

Out[ ]:

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight	Shell_weight
564	0	0.470	0.355	0.130	0.5465	0.2005	0.1260	
1584	0	0.515	0.375	0.110	0.6065	0.3005	0.1310	
3613	0	0.610	0.490	0.170	1.3475	0.7045	0.2500	
885	0	0.670	0.585	0.160	1.3090	0.5445	0.2945	
3106	1	0.300	0.220	0.065	0.1235	0.0590	0.0260	
...	...	...	...	...	...	...	...	...
56	2	0.445	0.350	0.120	0.4425	0.1920	0.0955	
1890	2	0.565	0.455	0.155	0.9355	0.4210	0.1830	
1299	1	0.530	0.415	0.110	0.5745	0.2525	0.1235	
3392	2	0.645	0.515	0.185	1.4605	0.5835	0.3155	
195	2	0.500	0.405	0.155	0.7720	0.3460	0.1535	

1045 rows × 8 columns



In [ ]:

```
y_test
```

Out[ ]:

	Age
564	15.5
1584	7.5
3613	12.5
885	11.5
3106	6.5
...	...
56	9.5
1890	12.5
1299	10.5
3392	20.5
195	13.5

1045 rows × 1 columns



In [ ]:

```
from sklearn.metrics import r2_score  
r2_score(mlr.predict(x_test),y_test)
```

Out[ ]:

0.0557067898421757